

Sections covered since Exam II: 10.1 to 10.9, 11.1 to 11.6, 12.1 to 12.5.3, 12.7, 13.1 to 13.3.3, 14.1 to 14.10.

1. Identify each part (A through F) of the slotted-page structure used to store variable-length records.

A	B	C	D	E	F
---	---	---	---	---	---

**A = # of entries and pointer to free space**

**B, C = pointers to records**

**D = free space**

**E, F = records**

2. For the B+ tree below,

a. What is  $n$ ?

**5**

b. How many values can be stored in non-leaf nodes?

**$\lceil n/2 \rceil$  to  $n$ , which is 3 to 5**

c. Would a new value fit into one of the existing leaf nodes?


**Yes, there's room in both leaf nodes, which can hold  $\lceil (n-1)/2 \rceil$  to  $n-1$  nodes, which is 2 to 4**

d. In which block would the value "Frederick" be stored?

**In block 2**

e. In which block would the value "Arbutus" be stored?

**In block 2**



3. Consider the following schedule.

Time	T1	T2
1	read(A)	
2	write(A)	
3		read(A)
4	read(B)	
5	write(B)	
6		write(A)

Is the schedule conflict serializable as  $\langle T1, T2 \rangle$ ? Explain why or why not.

**Yes. You can swap instructions at times 3 and 4 (they work on different objects). You can then swap instructions at times 4 and 5 (they work on different objects). This results in  $\langle T1, T2 \rangle$ .**

**Although it wasn't part of the question, note that it is NOT conflict serializable at  $\langle T2, T1 \rangle$ . To do this, you would need to move the instruction at time 3 up to time 1. However, you can't swap the instructions at time 2 and 3, because they affect the same object and at least one of them is a "write".**

4. Consider the following SQL query using R (A, B, C, D) and S (D, E), where r(R) and s(S). Consider the relational algebra expression, derived from the SQL. Using equivalence rules, suggest an equivalent relational algebra expression that is more efficient.

```
select  s.E, r.A
from    r, s
where   r.D = s.D
and     s.D = 'Baltimore'
and     r.B = 'Sales'
```

$$\Pi_{s.E, r.A} ( \sigma_{r.D = s.D \wedge s.D = 'Baltimore' \wedge r.B = 'Sales'} ( r \times s ) )$$

**There are a number of solutions. Consider our rules of thumb, and perform the select BEFORE the cross product.**

$$\Pi_{s.E, r.A} ( \sigma_{r.D = s.D} ( r.B = 'Sales' (r) \times s.D = 'Baltimore' (s) ) )$$

5. Given R (A, B, C, D) and S (D, E), where r(R) and s(S), assume that r has 1,000,000 rows with 100 rows stored per block, s has 100,000 rows with 500 rows stored per block, the block seek time is 0.4 microseconds, and the block transfer time is 0.1 microseconds, there is a primary index on A and a secondary index on B. Assume the height of any index used is 5. Assume there are 5 rows where B = 'Baltimore'. How long will it take to execute the following statements?

$\sigma_{B = \text{'Baltimore'}}(r)$

**You need algorithm A5 for secondary index, equality on non-key.**

$$\begin{aligned} & (\text{index height} + \text{number of rows}) * (\text{seek time} + \text{transfer time}) \\ & (5 + 5) * (0.4 + 0.1) * 10^{-6} \text{ seconds} \\ & = 5 * 10^{-6} \text{ seconds} \end{aligned}$$

$r \bowtie s$

**The Index-Nested-Loop join won't work, since there's no index on the join attribute (D). This leaves us with Nest-Loop or the Block-Nested-Loop joins.**

**Here's the answer for the Block-Nested-Loop join.**

**To retrieve data from r, we need --  $b_r * t_s + b_r * t_T$**

**To retrieve data from s, we need --  $b_r * t_s + b_r * b_s * t_T$**

**Note that there are 10,000 blocks in r and 200 blocks in s.**

$$\begin{aligned} & 10,000 * 0.4 * 10^{-6} + 10,000 * 0.1 * 10^{-6} + 10,000 * 0.4 * 10^{-6} + 10,000 * 200 * 0.1 * 10^{-6} \\ & = 0.209 \text{ seconds} \end{aligned}$$

**Out of curiosity, is it faster than the Nested-Loop join? Let's see ...**

**To retrieve data from r, we need --  $b_r * t_s + b_r * t_T$**

**To retrieve data from s, we need --  $n_r * t_s + n_r * b_s * t_T$**

**Note that there are 10,000 blocks in r and 200 blocks in s.**

**Note that the Nested-Loop join uses  $n_r$  as a multiplier.**

$$\begin{aligned} & 10,000 * 0.4 * 10^{-6} + 10,000 * 0.1 * 10^{-6} + 1,000,000 * 0.4 * 10^{-6} + 1,000,000 * 200 * 0.1 * 10^{-6} \\ & = 20.405 \text{ seconds} \end{aligned}$$